

1 Building A Pricing Model For First Time Home Buyers



2 Overview

In this analysis, I inspect the King County House Sales dataset and iteratively develop a multiple regression model to analyze house prices.

3 Business Problem

There are lots of residents in King County, Washington who are considering buying their first home. These prospective buyers could benefit immensely from being able to accurately forecast the price of their first home based on a set of given parameters.

As an analyst for DLG Real Estate Agency, I am tasked with developing a regression model to help my fellow employees determine which homes are best for their clients.

4 Importing Data, Necessary Libraries



```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.formula as smf
import statsmodels.api as sm
from statsmodels.formula.api import ols
from statsmodels.graphics.gofplots import qqplot
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
```

```
In [2]: import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: pd.set_option("display.max_columns", 100)
```

```
In [4]: df = pd.read_csv('data/kc_house_data.csv')
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   id                    21597 non-null  int64  
 1   date                  21597 non-null  object  
 2   price                 21597 non-null  float64 
 3   bedrooms              21597 non-null  int64  
 4   bathrooms             21597 non-null  float64 
 5   sqft_living           21597 non-null  int64  
 6   sqft_lot              21597 non-null  int64  
 7   floors                21597 non-null  float64 
 8   waterfront            19221 non-null  float64 
 9   view                  21534 non-null  float64 
10   condition             21597 non-null  int64  
11   grade                 21597 non-null  int64  
12   sqft_above            21597 non-null  int64  
13   sqft_basement         21597 non-null  object  
14   yr_built               21597 non-null  int64  
15   yr_renovated          17755 non-null  float64 
16   zipcode               21597 non-null  int64  
17   lat                   21597 non-null  float64 
18   long                  21597 non-null  float64 
19   sqft_living15         21597 non-null  int64  
20   sqft_lot15            21597 non-null  int64  
dtypes: float64(8), int64(11), object(2)
memory usage: 3.5+ MB
```

4.1 Column Names and descriptions for Kings County Data Set

- **id** - Unique identifier for a house
- **date** - Date house was sold
- **price** - Price is prediction target
- **bedrooms** - Number of Bedrooms/House
- **bathrooms** - Number of bathrooms/bedrooms
- **sqft_living** - Square footage of the home
- **sqft_lot** - Square footage of the lot
- **floors** - Total floors (levels) in house
- **waterfront** - House which has a view to a waterfront
- **view** - score of view from house
- **condition** - How good the condition is (Overall)
- **grade** - overall grade given to the housing unit, based on King County grading system
- **sqft_above** - square footage of house apart from basement
- **sqft_basement** - square footage of the basement
- **yr_built** - Built Year
- **yr_renovated** - Year when house was renovated
- **zipcode** - zip
- **lat** - Latitude coordinate
- **long** - Longitude coordinate
- **sqft_living15** - The square footage of interior housing living space for the nearest 15 neighbors
- **sqft_lot15** - The square footage of the land lots of the nearest 15 neighbors

Due to time constraints on this project, I am focusing solely on the following predictors:

- bedrooms
- bathrooms
- sqft_living
- sqft_lot
- floors
- waterfront
- condition
- grade
- yr_built
- zipcode
- view

Consequently, all other columns are dropped.

```
In [6]: to_drop = ['date', 'sqft_above', 'sqft_basement', 'yr_renovated', 'lat', 'l  
df.drop(to_drop, axis=1, inplace=True)
```

5 Basic Data Cleaning & Initial Model

Looking at the DataFrame information provided above, it appears that some columns have varying amounts of null values. Let's drop those and see if there are enough remaining entries for our analysis (at least 15,000).

```
In [7]: df.dropna(inplace=True)
```

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 19164 entries, 1 to 21596
Data columns (total 13 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   id              19164 non-null   int64
 1   price           19164 non-null   float64
 2   bedrooms        19164 non-null   int64
 3   bathrooms       19164 non-null   float64
 4   sqft_living     19164 non-null   int64
 5   sqft_lot        19164 non-null   int64
 6   floors          19164 non-null   float64
 7   waterfront      19164 non-null   float64
 8   view            19164 non-null   float64
 9   condition       19164 non-null   int64
10   grade           19164 non-null   int64
11   yr_built        19164 non-null   int64
12   zipcode         19164 non-null   int64
dtypes: float64(5), int64(8)
memory usage: 2.0 MB
```

```
In [9]: df.head()
```

Out[9]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
1	6414100192	538000.0	3	2.25	2570	7242	2.0	0.
2	5631500400	180000.0	2	1.00	770	10000	1.0	0.
3	2487200875	604000.0	4	3.00	1960	5000	1.0	0.
4	1954400510	510000.0	3	2.00	1680	8080	1.0	0.
5	7237550310	1230000.0	4	4.50	5420	101930	1.0	0.


I now create a DataFrame **df_pred** containing only our predictors, dropping the *price* & *id* columns.

```
In [10]: df_pred = df.drop(['price', 'id'], axis=1)
```

Next, I take a look at the correlation between these features:

```
In [11]: plt.figure(figsize=(12,9))
sns.heatmap(df_pred.corr())
```

Out[11]: <AxesSubplot:>



```
In [12]: corr_pair = df_pred.corr().abs().stack().reset_index().sort_values(0, ascen
corr_pair['pairs'] = list(zip(corr_pair.level_0, corr_pair.level_1))
corr_pair.set_index(['pairs'], inplace = True)
corr_pair.drop(columns=['level_1', 'level_0'], inplace = True)
corr_pair.columns = ['cc']
corr_pair.drop_duplicates(inplace=True)
```

```
In [13]: corr_pair[(corr_pair.cc>.75) & (corr_pair.cc <1)]
```

```
Out[13]:
```

cc	
pairs	
(sqft_living, grade)	0.763701
(bathrooms, sqft_living)	0.755909

As seen in both the heatmap & new DataFrame **corr_pair**, the variables *sqft_living*, *grade*, & *bathrooms* are highly correlated (correlation coefficient having an absolute value of over 0.75, indicated on the heatmap by a light shade).

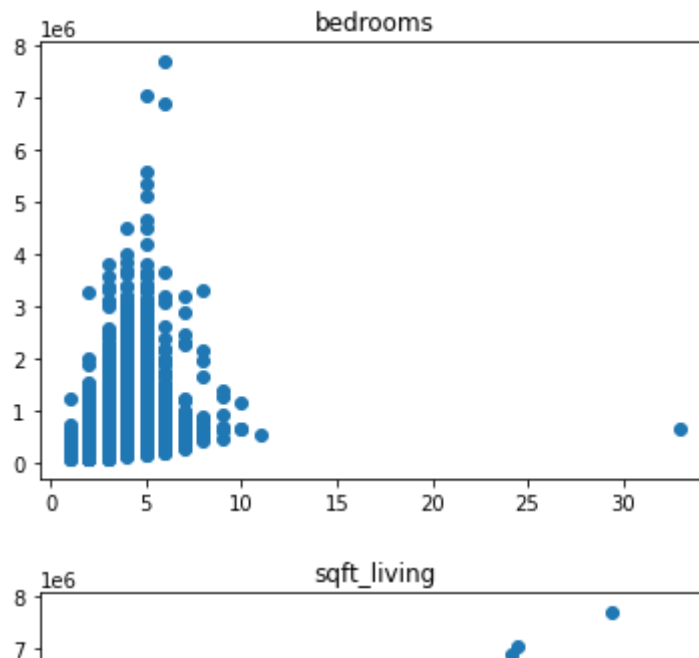
In order to remove collinear features, I drop *grade* & *bathrooms*, leaving only the *sqft_living* predictor.

```
In [14]: df_pred.drop(columns=['grade', 'bathrooms'], inplace=True)  
df.drop(columns=['grade', 'bathrooms'], inplace=True)
```

Next, I inspect how the remaining predictors look when plotted individually against the dependent *price* variable in a scatterplot.

These plots will be referenced again later on for any potential feature manipulation.

```
In [15]: for col in df_pred.columns:
          plt.scatter(df_pred[col], df['price'])
          plt.title(col)
          plt.show()
```



I now run a baseline regression model using the above set of predictors, unchanged, before evaluating which features to change.

```
In [16]: outcome = 'price'
          x_cols = df_pred.columns
          predictors = '+'.join(x_cols)

          f = outcome + '~' + predictors
```

```
In [17]: model_1 = ols(formula=f, data=df).fit()
model_1.summary()
```

Out[17]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.597
Model:	OLS	Adj. R-squared:	0.597
Method:	Least Squares	F-statistic:	3149.
Date:	Tue, 26 Jan 2021	Prob (F-statistic):	0.00
Time:	00:45:56	Log-Likelihood:	-2.6424e+05
No. Observations:	19164	AIC:	5.285e+05
Df Residuals:	19154	BIC:	5.286e+05
Df Model:	9		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	6.454e+06	3.51e+06	1.836	0.066	-4.35e+05	1.33e+07
bedrooms	-5.209e+04	2279.762	-22.848	0.000	-5.66e+04	-4.76e+04
sqft_living	306.9798	2.606	117.804	0.000	301.872	312.087
sqft_lot	-0.3727	0.043	-8.681	0.000	-0.457	-0.289
floors	7.485e+04	3811.423	19.638	0.000	6.74e+04	8.23e+04
waterfront	5.634e+05	2.15e+04	26.174	0.000	5.21e+05	6.06e+05
view	5.94e+04	2586.549	22.965	0.000	5.43e+04	6.45e+04
condition	1.905e+04	2861.467	6.656	0.000	1.34e+04	2.47e+04
yr_built	-2494.5937	75.864	-32.883	0.000	-2643.293	-2345.895
zipcode	-16.8060	35.267	-0.477	0.634	-85.932	52.320

Omnibus:	11751.006	Durbin-Watson:	1.977
Prob(Omnibus):	0.000	Jarque-Bera (JB):	433493.239
Skew:	2.378	Prob(JB):	0.00
Kurtosis:	25.809	Cond. No.	2.05e+08

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.05e+08. This might indicate that there are strong multicollinearity or other numerical problems.


```
In [18]: data = df.copy()

y = data['price']
X = data.drop(['price', 'id'], axis = 1)
```

```
In [19]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [20]: linreg = LinearRegression()
linreg.fit(X_train, y_train)

y_hat_train = linreg.predict(X_train)
y_hat_test = linreg.predict(X_test)
```

```
In [21]: mse_train = mean_squared_error(y_train, y_hat_train)
mse_test = mean_squared_error(y_test, y_hat_test)

print('Train MSE:', mse_train)
print('Test MSE:', mse_test)


print('RMSE Train:', np.sqrt(mse_train))
print('RMSE Test:', np.sqrt(mse_test))
```

```
Train MSE: 55834878369.075
Test MSE: 54133700042.9983
RMSE Train: 236294.0506425733
RMSE Test: 232666.49961478834
```

```
In [22]: residuals = (y_test - y_hat_test)

sm.qqplot(residuals, line = "r")
```

Out[22]:




This first model has an R-squared value of 0.597, and a Root Mean Square Error of around 240,000. Additionally, the residuals seem somewhat skewed based on the ends of the Q-Q Plot.

6 Model 2: Dropping Outliers, Using Heuristics & Z-Scores

Now that a baseline model has been constructed to improve on, let's return to the initial business problem: creating a price-prediction model for first time home buyers.

Recall the scatterplots constructed for the individual predictors against price. There are multiple features in the dataset with values that far exceed what would be found in a 'first home'.

```
In [23]: for col in df_pred.columns:
          plt.scatter(df_pred[col], df['price'])
          plt.title(col)
          plt.show()
```



6.1 Dropping Data Using Heuristics

First, take a look at the *bedrooms* scatterplot. There is an obvious outlier home with 33 bedrooms, which is certainly unfeasible for any first time buyer.

I begin by dropping this entry.

```
In [24]: df[df['bedrooms'] == 33]
```

Out[24]:

	id	price	bedrooms	sqft_living	sqft_lot	floors	waterfront	view	c
15856	2402100895	640000.0	33	1620	6000	1.0	0.0	0.0	

```
In [25]: df.drop(labels=15856, axis=0, inplace=True)
```

Now, I inspect the *bedrooms* column.

```
In [26]: df.sort_values(by=['bedrooms'], axis=0, ascending=False)
```


Out[26]:

	id	price	bedrooms	sqft_living	sqft_lot	floors	waterfront	view
8748	1773100755	520000.0	11	3000	4960	2.0	0.0	0.0
13301	627300145	1150000.0	10	4590	10920	1.0	0.0	2.0
19239	8812401450	660000.0	10	2920	3745	2.0	0.0	0.0
15147	5566100170	650000.0	10	3610	11914	2.0	0.0	0.0
4092	1997200215	599999.0	9	3830	6988	2.5	0.0	0.0
...
648	922049078	157000.0	1	870	26326	1.0	0.0	0.0
7368	7228501903	250000.0	1	780	1033	1.0	0.0	0.0
3380	8807900236	430000.0	1	630	1362	1.0	0.0	0.0
18261	2781600195	285000.0	1	1060	54846	1.0	1.0	4.0
17282	4047200825	400000.0	1	1390	60984	1.0	0.0	0.0

19163 rows × 11 columns

```
In [27]: df['bedrooms'].hist()
```

Out[27]: <AxesSubplot:>




Even if a potential first time buyer is a large family in need of more rooms than usual, it is difficult to envision such a client needing any more than 8 bedrooms.

```
In [28]: df = df[df['bedrooms'] < 9]
```

```
In [29]: df['bedrooms'].hist()
```


```
Out[29]: <AxesSubplot:>
```



6.2 Dropping Data Using Z-Scores of Continuous Variables

```
In [30]: df['price'].hist()
```

```
Out[30]: <AxesSubplot:>
```



```
In [31]: mean_price = df['price'].mean()
std3_price = 3*df['price'].std()
print(mean_price - std3_price, mean_price + std3_price)


-571334.4145640415 1653916.1387847904
```

From the histogram of home prices, we can tell that none of the prices fall more than 3 standard deviations below the mean price (given by the interval above). There are, however, outlying prices more than 3 standard deviations above the mean. I elect to filter these homes out of the dataset.

```
In [32]: upper_price = mean_price + std3_price  
df = df[df['price'] <= upper_price]
```

```
In [33]: df['price'].hist()
```


Out[33]: <AxesSubplot:>



Next, I look at the *sqft_living* predictor.

```
In [34]: df['sqft_living'].hist()
```

Out[34]: <AxesSubplot:>



```
In [35]: df['sqft_living'].describe()
```

```
Out[35]: count      18803.000000
         mean        2033.471999
         std         837.690802
         min         370.000000
         25%        1411.500000
         50%        1900.000000
         75%        2510.000000
         max         7480.000000
         Name: sqft_living, dtype: float64
```

```
In [36]: mean_living = df['sqft_living'].mean()
         std3_living = 3*df['sqft_living'].std()
         print(mean_living - std3_living, mean_living + std3_living)


-479.600405407758  4546.544403705902
```

The min of *sqft_living* does not fall under 3 standard deviations below the mean. But just by observing the max, it is apparent that at least one entry exceeds 3 standard deviations above the column's mean. I filter out any values exceeding this upper limit.

```
In [37]: upper_living = mean_living + std3_living
         df = df[df['sqft_living'] <= upper_living]
```

```
In [38]: df['sqft_living'].hist()
```


```
Out[38]: <AxesSubplot:>
```



Next, I look at the *sqft_lot* columns for any outliers.

```
In [39]: df['sqft_lot'].hist()
```

```
Out[39]: <AxesSubplot:>
```



```
In [40]: df['sqft_lot'].describe()
```

```
Out[40]: count      1.863200e+04
mean        1.442871e+04
std         3.869237e+04
min         5.200000e+02
25%         5.000000e+03
50%         7.520000e+03
75%         1.039525e+04
max         1.651359e+06
Name: sqft_lot, dtype: float64
```

```
In [41]: mean_lot = df['sqft_lot'].mean()
std3_lot = 3*df['sqft_lot'].std()
print(mean_lot - std3_lot, mean_lot + std3_lot)

-101648.40775254855  130505.82767526215
```


Just like with *sqft_living*, the min *sqft_lot* value isn't 3 standard deviations or more below the mean. On the flip side, though, at least the column's max is more than 3 standard deviations above the mean. Similarly to before, I filter out any values that exceed this upper limit.

```
In [42]: upper_lot = mean_lot + std3_lot
df = df[df['sqft_lot'] <= upper_lot]
```



```
In [43]: df['sqft_lot'].hist()
```

```
Out[43]: <AxesSubplot:>
```



Finally, to keep up to date with the dataset, let's see if we still have a sufficient amount of entries (at least 15,000).

```
In [44]: df.count()
```

```
Out[44]: id          18344
price          18344
bedrooms       18344
sqft_living    18344
sqft_lot       18344
floors         18344
waterfront     18344
view           18344
condition      18344
yr_built       18344
zipcode        18344
dtype: int64
```

Now, I update the **df_pred** DataFrame after making the above changes to the main DataFrame **df**. Once this is done, I'm ready to run the new regression model.

```
In [45]: df_pred = df.drop(['price', 'id'], axis=1)
```

```
In [46]: outcome = 'price'
x_cols = df_pred.columns
predictors = '+'.join(x_cols)

f = outcome + '~' + predictors
```

```
In [47]: model_2 = ols(formula=f, data=df).fit()
model_2.summary()
```

Out[47]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.510
Model:	OLS	Adj. R-squared:	0.510
Method:	Least Squares	F-statistic:	2120.
Date:	Tue, 26 Jan 2021	Prob (F-statistic):	0.00
Time:	00:46:04	Log-Likelihood:	-2.4779e+05
No. Observations:	18344	AIC:	4.956e+05
Df Residuals:	18334	BIC:	4.957e+05
Df Model:	9		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	8823.0658	2.72e+06	0.003	0.997	-5.33e+06	5.35e+06
bedrooms	-3.629e+04	1916.261	-18.935	0.000	-4e+04	-3.25e+04
sqft_living	237.4873	2.452	96.858	0.000	232.681	242.293
sqft_lot	-0.8521	0.107	-7.989	0.000	-1.061	-0.643
floors	7.65e+04	2991.009	25.575	0.000	7.06e+04	8.24e+04
waterfront	1.448e+05	2.16e+04	6.719	0.000	1.03e+05	1.87e+05
view	5.358e+04	2128.263	25.176	0.000	4.94e+04	5.78e+04
condition	2.018e+04	2208.811	9.136	0.000	1.58e+04	2.45e+04
yr_built	-1999.5005	59.079	-33.845	0.000	-2115.300	-1883.701
zipcode	39.7169	27.320	1.454	0.146	-13.833	93.266

Omnibus:	2808.019	Durbin-Watson:	1.968
Prob(Omnibus):	0.000	Jarque-Bera (JB):	7182.388
Skew:	0.856	Prob(JB):	0.00
Kurtosis:	5.543	Cond. No.	2.05e+08

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.05e+08. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [48]: data = df.copy()
```

```
y = data['price']  
X = data.drop(['price', 'id'], axis = 1)
```

```
In [49]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [50]: linreg = LinearRegression()  
linreg.fit(X_train, y_train)  
  
y_hat_train = linreg.predict(X_train)  
y_hat_test = linreg.predict(X_test)
```


```
In [51]: mse_train = mean_squared_error(y_train, y_hat_train)  
mse_test = mean_squared_error(y_test, y_hat_test)  
  
print('Train MSE:', mse_train)  
print('Test MSE:', mse_test)  
  
print('RMSE Train:', np.sqrt(mse_train))  
print('RMSE Test:', np.sqrt(mse_test))
```

```
Train MSE: 31747789192.635834  
Test MSE: 31235540437.570656  
RMSE Train: 178179.09302899663  
RMSE Test: 176735.79274603844
```

```
In [52]: residuals = (y_test - y_hat_test)

sm.qqplot(residuals, line = "r")
```

Out[52]:



Getting rid of the outliers did lower the model's R-squared value from 0.597 to 0.510. On the flip side, though, the RMSE did improve considerably (~175,000 here vs ~240,000 before). Additionally, the Q-Q Plot indicates that the residuals have become more Normally distributed, though they still have some right skew.


The *zipcode* feature does have a coefficient with a small t-statistic. This does not overly concern me, however, because I will soon get around to encoding this feature, which should eliminate the problem.

Ultimately, I choose to stick with the changes made here, in the hopes that transforming some of the features will sufficiently improve the R-squared value.

7 Model 3: Transforming Continuous Data

Let's look again at the scatterplots of individual predictors vs. home price:

```
In [53]: for col in df_pred.columns:
          plt.scatter(df_pred[col], df['price'])
          plt.title(col)
          plt.show()
```




Here, it is apparent that the two continuous variables in the set of features are the *sqft_living* & *sqft_lot* columns. Both columns' scatterplots have a "cloud-like" appearance with no apparent vertically-aligned clusters.

Next, I take another look at the histograms for the two continuous predictors to see how their distribution looks.

```
In [54]: df['sqft_living'].hist()
```

```
Out[54]: <AxesSubplot:>
```




The *sqft_living* values seem to be distributed fairly normally, with a slight right skew. I determine that the distribution is good enough as is and does not require any transformation.

Next, I inspect the *sqft_lot* distribution.

```
In [55]: df['sqft_lot'].hist()
```

```
Out[55]: <AxesSubplot:>
```




This predictor's distribution does not appear Normal at all. Additionally, from using the `.describe()` method on the column previously, I now know that all of its values are above 0. This makes the column a prime candidate for log-transformation.

```
In [56]: df['sqft_lot'] = np.log(df['sqft_lot'])
```

```
In [57]: df['sqft_lot'].hist()
```

```
Out[57]: <AxesSubplot:>
```



The distribution of the column's transformed values is much closer to a Normal one, and should lead to improvements in the next model iteration.

```
df['sqft_basement'] = (df['sqft_basement'] - mean_val) / basement_range
```

```
df['sqft_basement'].hist()
```

ignore cells above

```
In [58]: df_pred = df.drop(['price', 'id'], axis=1)
```

```
In [59]: outcome = 'price'
x_cols = df_pred.columns
predictors = '+'.join(x_cols)

f = outcome + '~' + predictors
```



```
In [60]: model_3 = ols(formula=f, data=df).fit()
model_3.summary()
```

Out[60]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.517
Model:	OLS	Adj. R-squared:	0.517
Method:	Least Squares	F-statistic:	2181.
Date:	Tue, 26 Jan 2021	Prob (F-statistic):	0.00
Time:	00:46:09	Log-Likelihood:	-2.4765e+05
No. Observations:	18344	AIC:	4.953e+05
Df Residuals:	18334	BIC:	4.954e+05
Df Model:	9		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.023e+07	2.77e+06	3.689	0.000	4.79e+06	1.57e+07
bedrooms	-3.558e+04	1894.968	-18.775	0.000	-3.93e+04	-3.19e+04
sqft_living	249.7833	2.540	98.357	0.000	244.806	254.761
sqft_lot	-3.752e+04	2057.473	-18.235	0.000	-4.16e+04	-3.35e+04
floors	5.565e+04	3229.210	17.234	0.000	4.93e+04	6.2e+04
waterfront	1.672e+05	2.14e+04	7.799	0.000	1.25e+05	2.09e+05
view	5.346e+04	2112.888	25.302	0.000	4.93e+04	5.76e+04
condition	1.968e+04	2193.067	8.973	0.000	1.54e+04	2.4e+04
yr_built	-2053.5894	58.737	-34.962	0.000	-2168.720	-1938.459
zipcode	-60.0071	27.781	-2.160	0.031	-114.461	-5.553

Omnibus:	2850.958	Durbin-Watson:	1.970
Prob(Omnibus):	0.000	Jarque-Bera (JB):	7557.925
Skew:	0.856	Prob(JB):	0.00
Kurtosis:	5.638	Cond. No.	2.09e+08

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.09e+08. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [61]: data = df.copy()
```

```
y = data['price']  
X = data.drop(['price', 'id'], axis = 1)
```

```
In [62]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [63]: linreg = LinearRegression()  
linreg.fit(X_train, y_train)  
  
y_hat_train = linreg.predict(X_train)  
y_hat_test = linreg.predict(X_test)
```


```
In [64]: mse_train = mean_squared_error(y_train, y_hat_train)  
mse_test = mean_squared_error(y_test, y_hat_test)  
  
print('Train MSE:', mse_train)  
print('Test MSE:', mse_test)  
  
print('RMSE Train:', np.sqrt(mse_train))  
print('RMSE Test:', np.sqrt(mse_test))
```

```
Train MSE: 30829252622.54272  
Test MSE: 32645058539.010094  
RMSE Train: 175582.60911190129  
RMSE Test: 180679.43584982242
```

```
In [65]: residuals = (y_test - y_hat_test)

sm.qqplot(residuals, line = "r")
```

Out[65]:



This latest model has multiple improvements, but they are relatively subtle.

The R-squared value increased slightly from 0.510 to 0.517. Additionally, none of the current features' coefficients have low t-scores.

8 Model 4: Dealing With Categorical Data

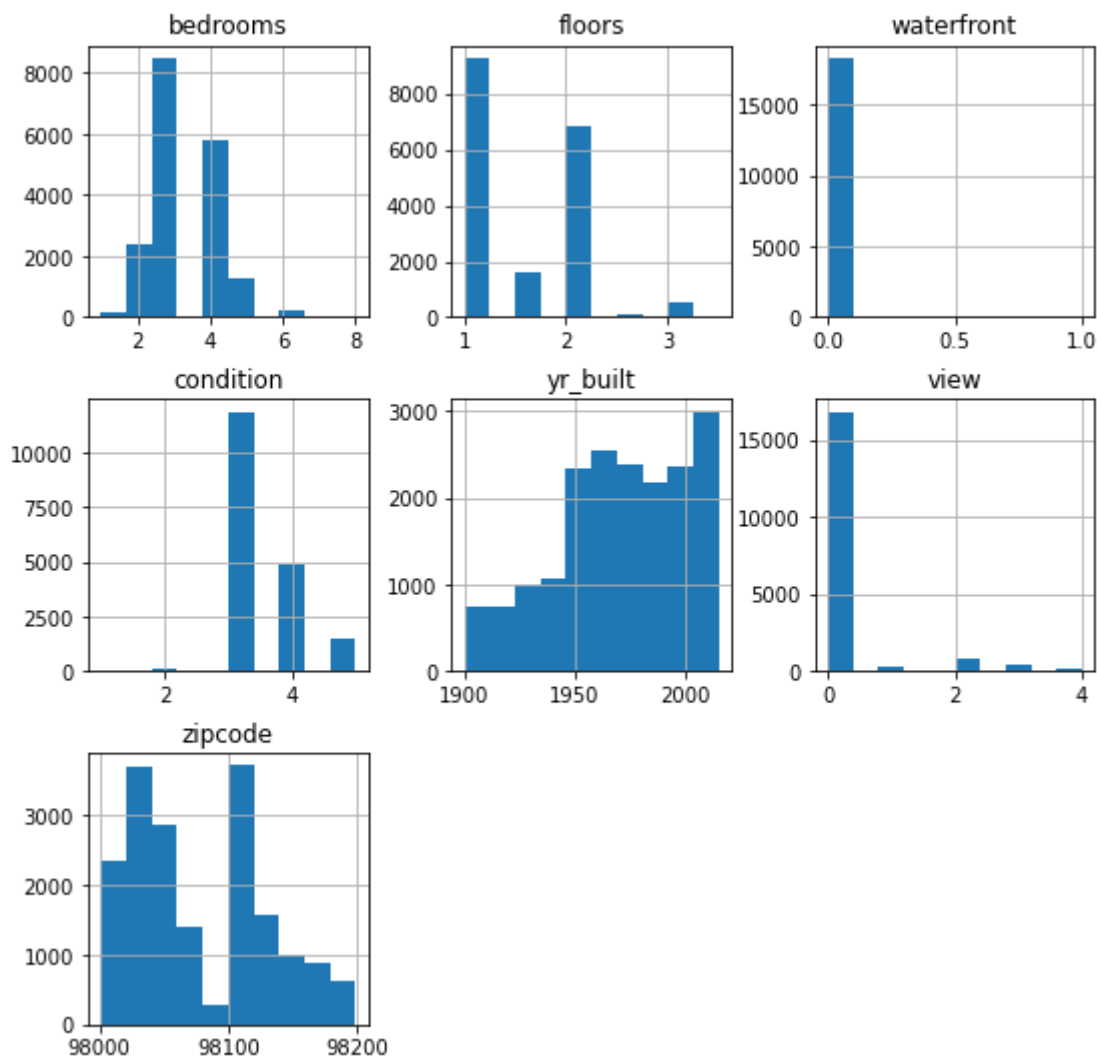
```
In [66]: cat = ['bedrooms', 'floors', 'waterfront', 'condition', 'yr_built', 'view',
```

```
In [67]: df[cat].info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 18344 entries, 1 to 21596
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   bedrooms    18344 non-null  int64
 1   floors      18344 non-null  float64
 2   waterfront  18344 non-null  float64
 3   condition   18344 non-null  int64
 4   yr_built    18344 non-null  int64
 5   view        18344 non-null  float64
 6   zipcode     18344 non-null  int64
dtypes: float64(3), int64(4)
memory usage: 1.1 MB
```

```
In [68]: df[cat].hist(figsize=(9,9))
```

```
Out[68]: array([[<AxesSubplot:title={'center':'bedrooms'}>,
<AxesSubplot:title={'center':'floors'}>,
<AxesSubplot:title={'center':'waterfront'}>],
[<AxesSubplot:title={'center':'condition'}>,
<AxesSubplot:title={'center':'yr_built'}>,
<AxesSubplot:title={'center':'view'}>],
[<AxesSubplot:title={'center':'zipcode'}>, <AxesSubplot:>,
<AxesSubplot:>]], dtype=object)
```




8.1 Binning The 'Year Built' Column Into Decades

```
In [69]: decade_bins = [1900, 1910, 1920, 1930, 1940, 1950, 1960, 1970, 1980, 1990,
```

```
In [70]: df['dec_built'] = pd.cut(df['yr_built'], bins=decade_bins,  
                                right=False)  
df['dec_built'] = df['dec_built'].cat.codes
```

```
In [71]: df['dec_built'].hist(figsize=(12,6))
```

Out[71]: <AxesSubplot:>



8.2 One Hot Encoding The 'Condition', 'Floors' & 'Zipcode' Columns

```
In [72]: cond_dummies = pd.get_dummies(df['condition'], prefix = 'cond', drop_first = True)  
floor_dummies = pd.get_dummies(df['floors'], prefix = 'floor', drop_first = True)  
zip_dummies = pd.get_dummies(df['zipcode'], prefix = 'zip', drop_first = True)
```

```
In [73]: df_d = pd.concat([df, cond_dummies, floor_dummies, zip_dummies], axis=1)
```

```
In [74]: df_d.columns = df_d.columns.str.replace('.', '_')  
df_d.drop(['condition', 'floors', 'zipcode', 'yr_built'], axis=1, inplace=True)
```

```
In [75]: df_d.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 18344 entries, 1 to 21596
Data columns (total 86 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    18344 non-null  int64
1   price                 18344 non-null  float64
2   bedrooms              18344 non-null  int64
3   sqft_living           18344 non-null  int64
4   sqft_lot              18344 non-null  float64
5   waterfront            18344 non-null  float64
6   view                  18344 non-null  float64
7   dec_built             18344 non-null  int8
8   cond_2                18344 non-null  uint8
9   cond_3                18344 non-null  uint8
10  cond_4                18344 non-null  uint8
11  cond_5                18344 non-null  uint8
12  floor_1_5            18344 non-null  uint8
13  floor_2_0            18344 non-null  uint8
14  floor_2_5            18344 non-null  uint8
15  floor_3_0            18344 non-null  uint8
16  floor_3_5            18344 non-null  uint8
17  zip_98002             18344 non-null  uint8
18  zip_98003             18344 non-null  uint8
19  zip_98004             18344 non-null  uint8
20  zip_98005             18344 non-null  uint8
21  zip_98006             18344 non-null  uint8
22  zip_98007             18344 non-null  uint8
23  zip_98008             18344 non-null  uint8
24  zip_98010             18344 non-null  uint8
25  zip_98011             18344 non-null  uint8
26  zip_98014             18344 non-null  uint8
27  zip_98019             18344 non-null  uint8
28  zip_98022             18344 non-null  uint8
29  zip_98023             18344 non-null  uint8
30  zip_98024             18344 non-null  uint8
31  zip_98027             18344 non-null  uint8
32  zip_98028             18344 non-null  uint8
33  zip_98029             18344 non-null  uint8
34  zip_98030             18344 non-null  uint8
35  zip_98031             18344 non-null  uint8
36  zip_98032             18344 non-null  uint8
37  zip_98033             18344 non-null  uint8
38  zip_98034             18344 non-null  uint8
39  zip_98038             18344 non-null  uint8
40  zip_98039             18344 non-null  uint8
41  zip_98040             18344 non-null  uint8
42  zip_98042             18344 non-null  uint8
43  zip_98045             18344 non-null  uint8
44  zip_98052             18344 non-null  uint8
45  zip_98053             18344 non-null  uint8
46  zip_98055             18344 non-null  uint8
47  zip_98056             18344 non-null  uint8
48  zip_98058             18344 non-null  uint8
49  zip_98059             18344 non-null  uint8
```



```
50 zip_98065      18344 non-null uint8
51 zip_98070      18344 non-null uint8
52 zip_98072      18344 non-null uint8
53 zip_98074      18344 non-null uint8
54 zip_98075      18344 non-null uint8
55 zip_98077      18344 non-null uint8
56 zip_98092      18344 non-null uint8
57 zip_98102      18344 non-null uint8
58 zip_98103      18344 non-null uint8
59 zip_98105      18344 non-null uint8
60 zip_98106      18344 non-null uint8
61 zip_98107      18344 non-null uint8
62 zip_98108      18344 non-null uint8
63 zip_98109      18344 non-null uint8
64 zip_98112      18344 non-null uint8
65 zip_98115      18344 non-null uint8
66 zip_98116      18344 non-null uint8
67 zip_98117      18344 non-null uint8
68 zip_98118      18344 non-null uint8
69 zip_98119      18344 non-null uint8
70 zip_98122      18344 non-null uint8
71 zip_98125      18344 non-null uint8
72 zip_98126      18344 non-null uint8
73 zip_98133      18344 non-null uint8
74 zip_98136      18344 non-null uint8
75 zip_98144      18344 non-null uint8
76 zip_98146      18344 non-null uint8
77 zip_98148      18344 non-null uint8
78 zip_98155      18344 non-null uint8
79 zip_98166      18344 non-null uint8
80 zip_98168      18344 non-null uint8
81 zip_98177      18344 non-null uint8
82 zip_98178      18344 non-null uint8
83 zip_98188      18344 non-null uint8
84 zip_98198      18344 non-null uint8
85 zip_98199      18344 non-null uint8
dtypes: float64(4), int64(3), int8(1), uint8(78)
memory usage: 2.5 MB
```

```
In [76]: df_pred = df_d.drop(['price', 'id'], axis=1)
```

```
In [77]: outcome = 'price'
x_cols = df_pred.columns
predictors = '+'.join(x_cols)

f = outcome + '~' + predictors
```

```
In [78]: model_4 = ols(formula=f, data=df_d).fit()
model_4.summary()
```

Out[78]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.808
Model:	OLS	Adj. R-squared:	0.807
Method:	Least Squares	F-statistic:	913.3
Date:	Tue, 26 Jan 2021	Prob (F-statistic):	0.00
Time:	00:46:13	Log-Likelihood:	-2.3921e+05
No. Observations:	18344	AIC:	4.786e+05
Df Residuals:	18259	BIC:	4.792e+05
Df Model:	84		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
--	------	---------	---	------	--------	--------

The sizable jump in R-squared value is very encouraging, but before I proceed with the rest of modeling, I observe the t-scores (and their associated p-values) for the predictor coefficients.

A few of the p-values for dummy variable coefficients exceed what most would consider an acceptable cutoff of $p=0.05$. Consequently, I drop these dummy columns before re-running the model.

```
In [79]: high_t_score = ['cond_2', 'floor_3_5', 'zip_98002', 'zip_98003', 'zip_98022', 'z
df_d.drop(high_t_score, axis=1, inplace=True)
```

```
In [80]: df_pred = df_d.drop(['price', 'id'], axis=1)
```

```
In [81]: outcome = 'price'
x_cols = df_pred.columns
predictors = '+'.join(x_cols)

f = outcome + '~' + predictors
```

```
In [82]: model_4 = ols(formula=f, data=df_d).fit()
model_4.summary()
```

Out[82]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.808
Model:	OLS	Adj. R-squared:	0.807
Method:	Least Squares	F-statistic:	1022.
Date:	Tue, 26 Jan 2021	Prob (F-statistic):	0.00
Time:	00:46:15	Log-Likelihood:	-2.3921e+05
No. Observations:	18344	AIC:	4.786e+05
Df Residuals:	18268	BIC:	4.792e+05
Df Model:	75		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
--	------	---------	---	------	--------	--------

A couple more zipcode dummy columns have coefficients with p-values over 0.05. I drop these last columns before finishing my model.

```
In [83]: df_d.drop(['zip_98042', 'zip_98070'], axis=1, inplace=True)
```

```
In [84]: df_pred = df_d.drop(['price', 'id'], axis=1)
```

```
In [85]: outcome = 'price'
x_cols = df_pred.columns
predictors = '+'.join(x_cols)

f = outcome + '~' + predictors
```

```
In [86]: model_4 = ols(formula=f, data=df_d).fit()
model_4.summary()
```

Out[86]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.808
Model:	OLS	Adj. R-squared:	0.807
Method:	Least Squares	F-statistic:	1050.
Date:	Tue, 26 Jan 2021	Prob (F-statistic):	0.00
Time:	00:46:17	Log-Likelihood:	-2.3922e+05
No. Observations:	18344	AIC:	4.786e+05
Df Residuals:	18270	BIC:	4.792e+05
Df Model:	73		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
--	------	---------	---	------	--------	--------

```
In [87]: data = df_d.copy()

y = data['price']
X = data.drop(['price', 'id'], axis = 1)
```

```
In [88]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [89]: linreg = LinearRegression()
linreg.fit(X_train, y_train)

y_hat_train = linreg.predict(X_train)
y_hat_test = linreg.predict(X_test)
```

```
In [90]: mse_train = mean_squared_error(y_train, y_hat_train)
mse_test = mean_squared_error(y_test, y_hat_test)

print('Train MSE:', mse_train)
print('Test MSE:', mse_test)



print('RMSE Train:', np.sqrt(mse_train))
print('RMSE Test:', np.sqrt(mse_test))
```

```
Train MSE: 12298396429.302235
Test MSE: 13033700783.34583
RMSE Train: 110898.13537342382
RMSE Test: 114165.23456528188
```

```
In [91]: residuals = (y_test - y_hat_test)

sm.qqplot(residuals, line = "r")
```

Out[91]:



9 Final Model Evaluation

My final regression model, built to predict the price of a house in King County, can be evaluated using the following metrics.

- **R-squared = 0.808** : This value indicates that my model explains 80.8% of the variation in home prices from their mean value.

- **Root Mean Square Error (Using Test Data) = 114522** : On average, my model's predicted price is +/- \$114,522 from the home's actual value.
- **Q-Q Plot (see above)** : From the plot, it seems that the residuals produced by my model have a relatively Normal distribution within 2 standard deviations of the mean. However, there is some skew, especially towards the right. In other words, my model is generally best at predicting the prices of homes that cost up to about \$1,009,014.

10 Conclusions

My analysis leads to the following advice for any prospective first time home buyer in King County, WA:

- Clients looking to save on their home purchase could start by looking in these zip codes: 98198, 98188, 98031, 98038, 98178, 98168 & 98058.
- Conversely, clients looking to make a bigger investment could start by looking in these zip codes: 98039, 98004, 98119, 98112, 98109, 98102 & 98040.
- Clients looking to save money should consider the home's condition grade, as it seems to have a sizable impact on price. Even going from an average grade to a high grade can increase a home's price significantly.
- As expected, there seems to be significant correlation between a home's square footage & its price.

11 Future Work

Given more time, I would look at the features I had to initially omit. This would involve using the *lat* & *long* columns to further inspect the impact of specific locations (beyond zip codes) on price. Additionally, I would look at the *sqft_living15* & *sqft_lot15* columns to get insight on what neighborhoods are most expensive to live in and the overall impact of comparative size of neighbors' homes.

In []: